PROJECT REPORT

Big Data ETL Pipeline with Machine Learning

for Network Intrusion Detection using Apache Spark

CICIDS2017 — Canadian Institute for Cybersecurity

Abstract

This report presents the design, implementation, and evaluation of a Big Data ETL (Extract, Transform, Load) pipeline integrated with a Machine Learning model for network intrusion detection. The pipeline was built using Apache Spark (PySpark 3.5+) on the CICIDS2017 Network Intrusion Detection dataset — a benchmark dataset containing approximately 2.8 million network flow records across 8 CSV files totalling ~2.8GB.

The pipeline extracts raw CSV data, applies comprehensive data cleaning (handling missing values, infinity values, duplicates, and outliers), performs feature engineering and standardization, loads the transformed data into Parquet and MySQL storage, and finally trains a Random Forest classifier to detect network attacks. The system was optimized for execution on a laptop with an Intel i3 11th Generation processor and 16GB RAM, achieving a ~43% performance improvement over unoptimized defaults. The pipeline follows a modular design pattern and is documented across etl_functions.py, ml_functions.py, etl_pipeline.py, and ml_pipeline.py.

Table of Contents

1. Project Introduction and Planning 3
2. Extract Phase — incl. Dataset Features (All 79 Columns) 4
3. Transform Phase 8
4. Load Phase 10
5. ETL Pipeline Implementation 11
6. ETL Pipeline Optimization 12
7. Machine Learning Integration 13
8. Model Evaluation 15
9. Challenges and Resolutions 16
10. Results and Conclusion 17

1. Project Introduction and Planning

1.1 Project Title

Big Data ETL Pipeline with Machine Learning for Network Intrusion Detection using Apache Spark

1.2 Objectives

- Build a modular, production-grade ETL pipeline using Apache Spark (PySpark)
- Process the CICIDS2017 dataset: ~2.8GB, 8 CSV files, 2.8 million rows, 79 features
- Clean, transform, and engineer features from raw network traffic data
- Load transformed data into Parquet files and a MySQL database
- Integrate a Random Forest machine learning model for intrusion detection
- Optimize the entire pipeline for a local machine (i3 11th Gen, 16GB RAM)
- Document all phases and present results

1.3 Dataset Overview

The CICIDS2017 (Canadian Institute for Cybersecurity Intrusion Detection System 2017) dataset is a widely used benchmark for network intrusion detection research. It contains realistic network traffic captured over 5 working days, with both benign and attack traffic clearly labelled.

Property	Details
Provider	Canadian Institute for Cybersecurity, University of New Brunswick
Total Size	~2.8 GB
Total Records	2,830,743 network flow rows
Features	79 numeric features per row
Label Column	BENIGN or specific attack type (14 classes)
Attack Categories	DoS, DDoS, PortScan, Brute Force, Web Attacks, Botnet, Infiltration
Download	https://www.unb.ca/cic/datasets/ids-2017.html

1.4 Dataset Files

CSV File	Attack Types
Monday-WorkingHours.pcap_ISCX.csv	Benign traffic only
Tuesday-WorkingHours.pcap_ISCX.csv	FTP-BruteForce, SSH-BruteForce
Wednesday-WorkingHours.pcap_ISCX.csv	DoS Hulk, DoS GoldenEye, Heartbleed
Thursday-Morning-WebAttacks.pcap_ISCX.csv	SQL Injection, XSS, Brute Force
Thursday-Afternoon-Infilteration.pcap_ISCX.csv	Infiltration
Friday-Morning.pcap_ISCX.csv	Botnet ARES
Friday-Afternoon-PortScan.pcap_ISCX.csv	PortScan
Friday-Afternoon-DDos.pcap_ISCX.csv	DDoS LOIT

1.5 Project Architecture

Pipeline Flow

Extract (8 CSVs) → Transform (clean + engineer + scale) → Load (Parquet + MySQL)
↓
ML Pipeline: Load Parquet → Prepare Features → Train RF Model → Evaluate → Predict

Files: etl_functions.py | ml_functions.py | etl_pipeline.py | ml_pipeline.py

2. Extract Phase

2.1 Overview

The extract phase reads all 8 CICIDS2017 CSV files from the dataset directory in a single Spark read operation. Spark automatically parallelises reading across files. A source_file column is added to track which CSV each row came from, enabling per-file analysis later.

2.2 Implementation

```
df = spark.read \
    .option("header", "true") \
    .option("mode", "DROPMALFORMED") \
    .schema(schema) \
    .csv("dataset/*.csv")
```

2.3 Extract Results

Metric	Value
Raw rows extracted	2,830,743
Columns extracted	80 (79 features + source_file)
Files read	8 CSV files in parallel
Mode	DROPMALFORMED — bad rows silently skipped
Column name issue	All columns had leading whitespace — stripped immediately

2.4 Dataset Features (All 79 Columns)

The CICIDS2017 dataset contains 79 network traffic features derived from network flow analysis. Each row represents one network flow (a connection between source and destination). Below is the complete schema grouped by feature category.

Group 1 — Packet Count Features

Column Name	Type	Description
Destination Port	Integer	Destination port number of the flow
Flow Duration	Integer	Duration of the flow in microseconds
Total Fwd Packets	Integer	Total packets in the forward direction
Total Backward Packets	Integer	Total packets in the backward direction
Total Length of Fwd Packets	Integer	Total size of packets in forward direction (bytes)
Total Length of Bwd Packets	Integer	Total size of packets in backward direction (bytes)

Group 2 — Packet Length Features

Column Name	Type	Description
Fwd Packet Length Max	Integer	Maximum size of forward packet
Fwd Packet Length Min	Integer	Minimum size of forward packet
Fwd Packet Length Mean	Double	Mean size of forward packets
Fwd Packet Length Std	Double	Standard deviation of forward packet sizes
Bwd Packet Length Max	Integer	Maximum size of backward packet
Bwd Packet Length Min	Integer	Minimum size of backward packet
Bwd Packet Length Mean	Double	Mean size of backward packets
Bwd Packet Length Std	Double	Standard deviation of backward packet sizes
Min Packet Length	Integer	Minimum length of a packet in the flow
Max Packet Length	Integer	Maximum length of a packet in the flow
Packet Length Mean	Double	Mean length of a packet in the flow
Packet Length Std	Double	Standard deviation of packet lengths
Packet Length Variance	Double	Variance of packet lengths in the flow

Group 3 — Flow Rate Features

Column Name	Type	Description
Flow Bytes/s	Double	Number of bytes per second in the flow (can be Inf)
Flow Packets/s	Double	Number of packets per second in the flow
Fwd Packets/s	Double	Number of forward packets per second
Bwd Packets/s	Double	Number of backward packets per second
Down/Up Ratio	Integer	Download to upload ratio of the flow

Group 4 — Inter-Arrival Time (IAT) Features

Column Name	Type	Description

Flow IAT Mean	Double	Mean time between two packets in the flow
Flow IAT Std	Double	Standard deviation of time between packets
Flow IAT Max	Integer	Maximum time between two packets in the flow
Flow IAT Min	Integer	Minimum time between two packets in the flow
Fwd IAT Total	Integer	Total time between two forward packets
Fwd IAT Mean	Double	Mean time between two forward packets
Fwd IAT Std	Double	Standard deviation of forward IAT
Fwd IAT Max	Integer	Maximum time between two forward packets
Fwd IAT Min	Integer	Minimum time between two forward packets
Bwd IAT Total	Integer	Total time between two backward packets
Bwd IAT Mean	Double	Mean time between two backward packets
Bwd IAT Std	Double	Standard deviation of backward IAT
Bwd IAT Max	Integer	Maximum time between two backward packets
Bwd IAT Min	Integer	Minimum time between two backward packets

Group 5 — TCP Flag Features

Column Name	Type	Description
Fwd PSH Flags	Integer	Number of PSH flags in forward direction
Bwd PSH Flags	Integer	Number of PSH flags in backward direction
Fwd URG Flags	Integer	Number of URG flags in forward direction
Bwd URG Flags	Integer	Number of URG flags in backward direction
FIN Flag Count	Integer	Number of FIN flags in the flow
SYN Flag Count	Integer	Number of SYN flags in the flow
RST Flag Count	Integer	Number of RST flags in the flow
PSH Flag Count	Integer	Number of PSH flags in the flow
ACK Flag Count	Integer	Number of ACK flags in the flow
URG Flag Count	Integer	Number of URG flags in the flow
CWE Flag Count	Integer	Number of CWE flags in the flow
ECE Flag Count	Integer	Number of ECE flags in the flow

Group 6 — Header and Segment Features

Column Name	Type	Description
Fwd Header Length	Long	Total bytes used for headers in forward direction
Fwd Header Length2	Long	Duplicate of Fwd Header Length (renamed from Fwd Header Length34)
Bwd Header Length	Integer	Total bytes used for headers in backward direction

Average Packet Size	Double	Mean size of a packet in the flow
Avg Fwd Segment Size	Double	Mean size of forward segments
Avg Bwd Segment Size	Double	Mean size of backward segments

Group 7 — Bulk Transfer Features

Column Name	Type	Description
Fwd Avg Bytes/Bulk	Integer	Average bytes bulk rate in forward direction
Fwd Avg Packets/Bulk	Integer	Average packets bulk rate in forward direction
Fwd Avg Bulk Rate	Integer	Average bulk rate in the forward direction
Bwd Avg Bytes/Bulk	Integer	Average bytes bulk rate in backward direction
Bwd Avg Packets/Bulk	Integer	Average packets bulk rate in backward direction
Bwd Avg Bulk Rate	Integer	Average bulk rate in the backward direction

Group 8 — Subflow and Window Features

Column Name	Type	Description
Subflow Fwd Packets	Integer	Average packets in a forward subflow
Subflow Fwd Bytes	Integer	Average bytes in a forward subflow
Subflow Bwd Packets	Integer	Average packets in a backward subflow
Subflow Bwd Bytes	Integer	Average bytes in a backward subflow
Init_Win_bytes_forward	Integer	Total bytes sent in initial window — forward
Init_Win_bytes_backward	Integer	Total bytes sent in initial window — backward
act_data_pkt_fwd	Integer	Count of packets with at least 1 byte of data in forward direction
min_seg_size_forward	Integer	Minimum segment size in the forward direction

Group 9 — Active / Idle Time Features

Column Name	Type	Description
Active Mean	Double	Mean time a flow was active before becoming idle
Active Std	Double	Standard deviation of active time
Active Max	Integer	Maximum time a flow was active
Active Min	Integer	Minimum time a flow was active
Idle Mean	Double	Mean time a flow was idle before becoming active
Idle Std	Double	Standard deviation of idle time
Idle Max	Integer	Maximum time a flow was idle
Idle Min	Integer	Minimum time a flow was idle

Group 10 — Label and Engineered Columns

Column Name	Type	Description
Label	String	Original class label: BENIGN or attack type name
is_attack	Integer	Binary label: 0 = BENIGN, 1 = any attack (added by pipeline)
label_index	Double	Multi-class integer index from StringIndexer (added by pipeline)
source_file	String	Name of the CSV file the row was extracted from
pkt_byte_ratio	Double	Engineered: Total Fwd Packets / Total Length of Fwd Packets
fwd_bwd_ratio	Double	Engineered: Total Fwd Packets / Total Backward Packets
duration_bucket	String	Engineered: Categorical bucket of Flow Duration (very_short/short/medium/long)

3. Transform Phase

3.1 Overview

The transform phase consists of 8 sequential steps implemented as standalone functions in etl_functions.py. Each function is responsible for exactly one transformation, making the pipeline easy to test, debug and maintain independently.

3.2 Step-by-Step Transformations

Step 3.1 — Column Name Cleaning

CICIDS2017 has two columns both named 'Fwd Header Length'. Spark internally names the duplicate 'Fwd Header Length34'. This causes an AnalysisException (dot treated as struct accessor) unless renamed.

```
df = df.toDF(*[c.strip() for c in df.columns])
df = df.withColumnRenamed("Fwd Header Length.1", "Fwd Header Length2")
```

Step 3.2 — Missing Values

Rows with a null Flow Duration (the primary flow identifier) are dropped. All remaining numeric nulls are filled with 0 to preserve row count.

```
# Fill nulls  with 0
df = df.fillna(0, subset=["Flow Bytes/s", "Flow Packets/s"])
```

Step 3.3 — Infinity Values

Flow Bytes/s and Flow Packets/s are calculated as Total Bytes / Flow Duration. When Flow Duration = 0, the result is Infinity. A null check revealed 1,358 infinity values in Flow Bytes/s which were replaced with 0.

```
# — 1. Replace Infinity values first ————————————————————————
inf_cols = ["Flow Bytes/s", "Flow Packets/s"]

for c in inf_cols:
    df = df.withColumn(c,
        when(col(c) == float("inf"),  None)
       .when(col(c) == float("-inf"), None)
       .otherwise(col(c))
    )
```

Step 3.4 — Duplicate Removal

Full row deduplication removes rows that are 100% identical across all columns. This correctly identified 308,381 true duplicate flow records.

```
#Check duplicate count
total = df.count()
distinct = df.dropDuplicates().count()
df = df.dropDuplicates()
```

Step 3.5 — Outlier Capping (IQR Method)

IQR-based capping is applied to Flow Duration. Values beyond Q1 - 1.5*IQR or Q3 + 1.5*IQR are capped at the boundary rather than dropped, preserving attack rows that may have extreme durations.

```
from pyspark.sql.functions import percentile_approx

q1, q3 = df.select(
    percentile_approx("Flow Duration", 0.25),
    percentile_approx("Flow Duration", 0.75)
).first()

iqr = q3 - q1
lower, upper = q1 - 1.5 * iqr, q3 + 1.5 * iqr

df = df.withColumn("Flow Duration",
    when(col("Flow Duration") < lower, lower)
    .when(col("Flow Duration") > upper, upper)
    .otherwise(col("Flow Duration"))
)
```

Step 3.6 — Label Encoding

Two label columns are created for different ML tasks: is_attack (binary) and label_index (multi-class integer via StringIndexer).

```
# Binary: attack or not
df = df.withColumn("is_attack",
    when(col("Label") == "BENIGN", 0).otherwise(1))

# Multi-class index
indexer = StringIndexer(inputCol="Label", outputCol="label_index")
df = indexer.fit(df).transform(df)
```

Step 3.7 — Feature Engineering

Three new derived features are added to capture relationships not present in the raw data:

Feature	Formula	Purpose
pkt_byte_ratio	Total Fwd Packets / Total Length of Fwd Packets	Packet size efficiency
fwd_bwd_ratio	Total Fwd Packets / Total Backward Packets	Traffic direction asymmetry

Step 3.8 — Standardization

All numeric feature columns are assembled into a single vector using VectorAssembler, then normalized to zero mean and unit variance using StandardScaler. This prevents features with large value ranges from dominating the ML model.

```
assembler = VectorAssembler(
    inputCols=feature_cols,
    outputCol="features_raw",
    handleInvalid="skip"
)
df_vec = assembler.transform(df)

scaler = StandardScaler(inputCol="features_raw", outputCol="features",
                        withMean=True, withStd=True)
df_scaled = scaler.fit(df_vec).transform(df_vec)
```

3.3 Transform Results

Metric	Value
Rows after cleaning	2,522,362
Duplicates removed	308,381
Infinity values fixed	1,358 in Flow Bytes/s
New features added	2 (pkt_byte_ratio, fwd_bwd_ratio)
Total columns after transform	85

4. Load Phase

4.1 Parquet Storage

The primary output format is Parquet with Snappy compression. Parquet is columnar, compressed, and supports predicate pushdown — Spark only reads the columns and partitions it needs when querying. Data is partitioned by is_attack so benign and attack rows can be read independently.

```
df_scaled.repartition(4) \
    .write \
    .mode("overwrite") \
    .partitionBy("is_attack") \
    .parquet("./output/cicids2017_cleaned/")
```

4.3 Load Results

Metric	Value
Total rows written	2,564,722
Benign rows (is_attack=0)	2,138,844 (83.4%)
Attack rows (is_attack=1)	425,878 (16.6%)
Output columns	85
Parquet partitions	8 files (4 per is_attack partition)
Compression	appy (3-5x smaller than CSV)

5. ETL Pipeline Implementation

5.1 Modular Design

The pipeline follows a modular design pattern where every transformation is a standalone function. The pipeline class simply orchestrates these functions in the correct sequence. This makes each step independently testable, reusable, and replaceable.

File	Contains	Purpose
etl_functions.py	read_csvs, clean, transform, load functions	All ETL standalone functions
ml_functions.py	prepare, train, evaluate, predict functions	All ML standalone functions
etl_pipeline.py	CICIDS_ETL_Pipeline class	Orchestrates ETL phases
ml_pipeline.py	CICIDS_ML_Pipeline class	Orchestrates ML phases

5.2 Pipeline Class

```
def transform(self, df):
    df = clean_column_names(df)
    df = handle_missing_values(df)
    df = replace_infinity(df)
    df = remove_duplicates(df)
    df = cap_outliers(df, column="Flow Duration")
    df = encode_labels(df)
    df = engineer_features(df)
    df = scale_features(df)
    return df
```

6. ETL Pipeline Optimization

6.1 Spark Configuration

Config	Value	Reason
spark.master	local[6]	6 of 8 threads — leave 2 for OS
spark.driver.memory	10g	Handles 1 GB dataset in memory
spark.memory.offHeap.size	2g	Reduces JVM garbage collection pauses
spark.memory.fraction	0.8	More RAM allocated to execution
shuffle.partitions	8	Default 200 is too many for a laptop
parquet.filterPushdown	true	Reads only required rows from Parquet
compression.codec	snappy	3-5x smaller Parquet output files
KryoSerializer	enabled	10x faster serialization than Java default
worker.cleanup.interval	300s	Auto-deletes shuffle files every 5 min

6.2 Performance Improvement

Phase	Before	After
CSV Read (8 files)	~2.5 min	~1.8 min
Transform pipeline	~6.0 min	~3.5 min
Write Parquet	~3.0 min	~1.8 min
Full ETL pipeline	~12 min	~7 min

7. Machine Learning Integration

7.1 Overview

The machine learning phase is implemented as a separate pipeline (ml_pipeline.py) that reads directly from the ETL Parquet output. This separation means the ETL and ML phases can run independently — the model can be retrained without re-running ETL, and predictions can be made without retraining.

Why Separate ML Pipeline?

ETL runs only when raw data changes. ML runs independently for training or prediction.
Easier to debug — ETL failures and ML failures are isolated.
Industry standard: ETL and ML are separately schedulable jobs in production.

7.2 Feature Preparation

Numeric columns are assembled into a feature vector. Class weights are calculated to handle the 83%/17% class imbalance — without this, the model would be biased toward predicting BENIGN.

```
# Calculate class weights to compensate imbalance
total    = df_ml.count()
benign   = df_ml.filter(col("is_attack") == 0).count()
attack   = df_ml.filter(col("is_attack") == 1).count()

weight_benign = total / (2.0 * benign)
weight_attack = total / (2.0 * attack)

# Add weight column
from pyspark.sql.functions import when
df_ml = df_ml.withColumn("classWeight",
    when(col("is_attack") == 0, weight_benign)
    .otherwise(weight_attack))
```

7.3 Train / Test Split

```
# 80% train, 20% test — stratified by is_attack
train_df, test_df = df_ml.randomSplit([0.8, 0.2], seed=42)

# Verify split is balanced
print(f"Train rows : {train_df.count():,}")
print(f"Test rows  : {test_df.count():,}")
```

7.4 Random Forest Classifier

Random Forest was selected as the primary algorithm because it handles high-dimensional data well, is robust to outliers, provides feature importance scores, and does not require feature scaling (though scaled features were available).

```
rf = RandomForestClassifier(
    featuresCol="features",
    labelCol="is_attack",
    weightCol="classWeight",
    numTrees=50,           # start with 50, tune later
    maxDepth=10,
    maxBins=32,
    seed=42
)
```

7.5 Hyperparameter Tuning

Cross validation with grid search was applied on a 20% sample of the training data (to stay within laptop

```
from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
from pyspark.ml.evaluation import BinaryClassificationEvaluator

# Use Random Forest
rf_tuning = RandomForestClassifier(
    featuresCol="features",
    labelCol="is_attack",
    weightCol="classWeight",
    seed=42
)

# Define parameter grid
paramGrid = ParamGridBuilder() \
    .addGrid(rf_tuning.numTrees,  [20, 50, 100]) \
    .addGrid(rf_tuning.maxDepth,  [5, 10, 15]) \
    .addGrid(rf_tuning.maxBins,   [16, 32]) \
    .build()

print(f"Total parameter combinations: {len(paramGrid)}")
# 3 x 3 x 2 = 18 combinations
```

memory limits). The parameter grid tested 18 combinations across numTrees, maxDepth, and maxBins.

8. Model Evaluation

8.1 Model Comparison

Model	Accuracy	Precision	Recall	F1	AUC
Logistic Regression	~0.91	~0.90	~0.91	~0.90	~0.95
Decision Tree	~0.96	~0.96	~0.96	~0.96	~0.96
Random Forest	~0.98	~0.98	~0.98	~0.98	~0.99
Random Forest (Tuned)	~0.99	~0.99	~0.99	~0.99	~0.999

Best Model: Random Forest (Tuned)

numTrees=100, maxDepth=15, maxBins=32 selected by 3-fold cross validation.
AUC-ROC ~0.999 means the model correctly distinguishes benign from attack traffic with near-perfect accuracy.
F1 Score ~0.99 confirms both precision and recall are balanced — few false positives and false negatives.

8.2 Class Imbalance Handling

Important Note on Class Imbalance

Dataset: 83.4% BENIGN vs 16.6% ATTACK — heavily imbalanced.
Without class weights, a naive model predicts BENIGN for everything and achieves 83% accuracy without learning anything.
Solution applied: classWeight column assigns higher weight (x3.01) to attack samples during training.
Evaluation metric: AUC-ROC and F1-score used instead of accuracy to correctly measure performance on imbalanced data.

9. Challenges and Resolutions

	Challenge	Cause	Resolution
	Out of Memory (OOM) Error	Attempted to cache full 2.8GB DataFrame — expanded to 8GB+ in memory	Removed .cache(); used Parquet disk checkpoint instead
	Storage disk full	Spark shuffle files accumulated in /tmp during long operations	Set worker.cleanup.interval=300s; added cleanup() function to pipeline
	1,358 nulls in Flow Bytes/s	Infinity values from zero-duration flows not handled before fillna	Added explicit infinity replacement step before fillna

10. Results and Conclusion

10.1 Final Pipeline Results

Metric	Value
Raw rows extracted	2,830,743
Rows after full cleaning	2,564,722
Rows removed	266,021 (duplicates + nulls + malformed)
Original columns	79 features + Label
Final columns	85 (+ 6 engineered/encoded columns)
BENIGN rows	2,138,844 (83.4%)
Attack rows	425,878 (16.6%)
Best model	Random Forest (Tuned)
Best F1 Score	~0.99
Best AUC-ROC	~0.999
Pipeline speedup	~43% faster after optimization

10.2 Output Structure

Final Output Files

output/cleaned/is_attack=0/ ← 2,138,844 benign rows (Parquet)
output/cleaned/is_attack=1/ ← 425,878 attack rows (Parquet)
output/cleaned_label_summary/ ← per-label count CSV
output/predictions/ ← model prediction results (Parquet)
output/models/random_forest/ ← saved Random Forest model
MySQL: cicids_db.network_traffic ← full cleaned dataset

10.3 Conclusion

This project successfully demonstrates a complete, production-grade Big Data ETL pipeline integrated with a Machine Learning model using Apache Spark. Starting from raw network traffic CSV files, the pipeline extracts, cleans, transforms, and loads 2.8 million records while handling real-world data quality issues including leading whitespace in column names, infinity values, class imbalance, duplicate records, and outliers.

The modular design pattern — separating transformation functions from pipeline orchestration — proved highly effective for debugging and maintenance. Eight challenges were encountered and resolved during development, each providing deeper insight into Spark internals and production ETL design.

The integrated Random Forest classifier achieved ~0.99 F1 Score and ~0.999 AUC-ROC, demonstrating that the cleaned and engineered feature set produces high-quality inputs for machine learning. The pipeline optimization techniques reduced total execution time by ~43%, making the full pipeline feasible on a laptop environment.

End of Report

Big Data ETL Pipeline with ML for Network Intrusion Detection